

Load the given Teen Mental Health Dataset and display the first 10 rows.

```
In [39]: import pandas as pd
teen_health = pd.read_csv(r'C:\Users\Omkaar\Desktop\Christ\4th_Trimester\ML\ML-Lab-2547237\datasets\teen_health.csv')
teen_health.head()
```

Out[39]:

	age	gender	daily_social_media_hours	platform_usage	sleep_hours	screen_time_before_sleep	academic_performance	physical_activity	social_interaction_level	stress_level	anxiety_level	addiction_level	depression_label
0	14.0	NaN	7.9	Instagram	7.4	2.9	3.01	1.5	low	2.0	NaN	1.0	0.0
1	19.0	female	1.9	TikTok	8.0	2.9	3.22	0.8	NaN	8.0	1.0	10.0	0.0
2	17.0	female	1.3	Instagram	7.6	0.5	3.92	0.0	high	2.0	4.0	2.0	0.0
3	15.0	male	7.4	TikTok	6.9	1.6	NaN	0.8	NaN	1.0	7.0	9.0	NaN
4	15.0	female	NaN	Both	4.9	3.0	2.37	1.4	NaN	3.0	5.0	2.0	0.0

List all numerical and categorical variables separately.

```
In [40]: numerical_cols = teen_health.select_dtypes(include=['number']).columns.tolist()
categorical_cols = teen_health.select_dtypes(include=['object', 'category']).columns.tolist()

print(f"Numerical Columns: {numerical_cols}")
print(f"Categorical Columns: {categorical_cols}")
```

Numerical Columns: ['age', 'daily_social_media_hours', 'sleep_hours', 'screen_time_before_sleep', 'academic_performance', 'physical_activity', 'stress_level', 'anxiety_level', 'addiction_level', 'depression_label']
Categorical Columns: ['gender', 'platform_usage', 'social_interaction_level']

Check the data types of all columns and interpret whether any column requires type conversion

```
In [41]: teen_health.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1200 entries, 0 to 1199
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0    age                 1105 non-null  float64
1    gender              972 non-null   object
2    daily_social_media_hours  1056 non-null  float64
3    platform_usage       1032 non-null  object
4    sleep_hours         1008 non-null  float64
5    screen_time_before_sleep  1140 non-null  float64
6    academic_performance   1080 non-null  float64
7    physical_activity     1092 non-null  float64
8    social_interaction_level  1032 non-null  object
9    stress_level         1092 non-null  float64
10   anxiety_level        1092 non-null  float64
11   addiction_level       1128 non-null  float64
12   depression_label      1008 non-null  float64
dtypes: float64(10), object(3)
memory usage: 122.0+ KB
```

Interpretation:

- 1. age is to be int not float
- 2. if physcial activity is counted as number of hours, then its needed to be rounded off to the nearest ceiling value for easier interpretation

2.a. Generate descriptive statistics for all numerical variables. Which variable shows the highest variation?

```
In [42]: teen_health[numerical_cols].describe()
```

Out[42]:

	age	daily_social_media_hours	sleep_hours	screen_time_before_sleep	academic_performance	physical_activity	stress_level	anxiety_level	addiction_level	depression_label
count	1105.000000	1056.000000	1008.000000	1140.000000	1080.000000	1092.000000	1092.000000	1092.000000	1128.000000	1008.000000
mean	15.901357	4.544129	6.443155	1.735965	2.986148	1.016392	5.454212	5.608974	5.572695	0.025794
std	2.016294	2.027443	1.457104	0.717721	0.578481	0.581912	2.898641	2.851348	2.836578	0.158598
min	13.000000	1.000000	4.000000	0.500000	2.000000	0.000000	1.000000	1.000000	1.000000	0.000000
25%	14.000000	2.800000	5.100000	1.100000	2.490000	0.500000	3.000000	3.000000	3.000000	0.000000
50%	16.000000	4.500000	6.400000	1.700000	2.980000	1.000000	5.000000	6.000000	6.000000	0.000000
75%	18.000000	6.300000	7.700000	2.325000	3.480000	1.500000	8.000000	8.000000	8.000000	0.000000
max	19.000000	8.000000	9.000000	3.000000	4.000000	2.000000	10.000000	10.000000	10.000000	1.000000

2.b. Compare the mean and median of daily social media hours and sleep hours. What does the difference indicate about the distribution and skewness of these variables?

```
In [43]: import math

mean_social_media_hrs = teen_health['daily_social_media_hours'].mean()
mean_sleep_hrs = teen_health['sleep_hours'].mean()
```

```
print("Comparison")
print(f"Social Media hours mean: {round(mean_social_media_hrs,2)}")
print(f"Sleep hours mean: {round(mean_sleep_hrs,2)}")
print(f"Difference: {round(abs(mean_social_media_hrs - mean_sleep_hrs),2)}")
```

Comparison
Social Media hours mean: 4.54
Sleep hours mean: 6.44
Difference: 1.9

Interpretation: For a healthy teenager, it is prescribed to sleep for atleast 7-8 hours, and due to such a low difference in between sleep and social media hours, it clearly states that over usage of social media lacks the time for proper sleep

3. a. Identify all variables containing missing values. Calculate the number and percentage of missing values in each column.

```
In [44]: # identifying columns and numbers
print("Number of missing values in each column")
teen_health.isnull().sum()
```

Number of missing values in each column

```
Out[44]: age          95
gender        228
daily_social_media_hours  144
platform_usage  168
sleep_hours   192
screen_time_before_sleep   60
academic_performance  120
physical_activity  108
social_interaction_level  168
stress_level    108
anxiety_level   108
addiction_level   72
depression_label  192
dtype: int64
```

```
In [45]: # percentage
print("Percentage of missing values in each column")
round(teen_health.isnull().mean() * 100,2)
```

Percentage of missing values in each column

```
Out[45]: age          7.92
gender        19.00
daily_social_media_hours  12.00
platform_usage  14.00
sleep_hours   16.00
screen_time_before_sleep   5.00
academic_performance  10.00
physical_activity   9.00
social_interaction_level  14.00
stress_level       9.00
anxiety_level       9.00
addiction_level      6.00
depression_label    16.00
dtype: float64
```

b. Suggest suitable imputation techniques for missing values in daily social media hours, sleep hours, and screen time before sleep.

For daily social media hours, some suggestions would be:

1. Mean/Median Imputation - its simple and quick
2. Regression-based - uses age/gender

For sleep hours:

1. Median Imputation - more robust than mean for sleep data
2. Group based median - by age/gender
3. Forward backward fill

USAGE: The SimpleImputer class in scikit-learn is a built-in tool used to fill in missing values (NaN or blanks) in a dataset. Most machine learning models cannot handle missing data, so SimpleImputer replaces those gaps with a calculated baseline like the average or the most frequent value

```
In [52]: # example for mean imputation on social media hours
import numpy as np
from sklearn.impute import SimpleImputer

imputer_mean = SimpleImputer(strategy='median')
teen_health['daily_social_media_hours'] = imputer_mean.fit_transform(teen_health[['daily_social_media_hours']])

print('After Computation, nulls in social media hours:', teen_health['daily_social_media_hours'].isnull().sum())
teen_health.head(10)
```

After Computation, nulls in social media hours: 0

Out[52]:

	age	gender	daily_social_media_hours	platform_usage	sleep_hours	screen_time_before_sleep	academic_performance	physical_activity	social_interaction_level	stress_level	anxiety_level	addiction_level	depression_label
0	14.0	NaN	7.900000	Instagram	7.4	2.9	3.01	1.5	low	2.0	NaN	1.0	0.0
1	19.0	female	1.900000	TikTok	8.0	2.9	3.22	0.8	NaN	8.0	1.0	10.0	0.0
2	17.0	female	1.300000	Instagram	7.6	0.5	3.92	0.0	high	2.0	4.0	2.0	0.0
3	15.0	male	7.400000	TikTok	6.9	1.6	NaN	0.8	NaN	1.0	7.0	9.0	NaN
4	15.0	female	4.544129	Both	4.9	3.0	2.37	1.4	NaN	3.0	5.0	2.0	0.0
5	NaN	female	7.400000	Both	4.4	2.4	2.63	0.6	high	3.0	5.0	7.0	0.0
6	18.0	female	2.500000	Instagram	6.4	2.4	2.63	0.7	low	2.0	2.0	5.0	0.0
7	16.0	male	4.000000	Both	4.2	0.5	2.40	NaN	low	6.0	10.0	NaN	0.0
8	19.0	female	3.300000	TikTok	5.0	2.1	NaN	0.9	high	1.0	10.0	9.0	0.0
9	15.0	male	1.900000	TikTok	4.9	1.5	3.77	1.1	high	1.0	1.0	4.0	0.0

c. Suggest suitable imputation techniques for missing values in stress level, anxiety level, and addiction level.

Stress Level

-
1. Mode Imputation (most frequent value - simple & interpretable)

2. Forward Fill/Backward Fill (if survey data)

3. Group-based Mode (by age/gender demographics)

Anxiety Level

-
1. Median Imputation (robust for ordinal scales)

2. Domain-based (correlate with stress & sleep)

Addiction Level

-
1. Correlation-based (correlate with screen_time_before_sleep & social_media_hours)

2. Group-based Median (by demographic groups)

d. Explain how missingness in mental health indicators such as anxiety level could potentially be MAR or MNAR.

MAR (Missing At Random):

-
- If anxiety level is missing for teens who reported high stress levels

Pattern: missingness depends on observed stress_level, but NOT on unobserved anxiety

MNAR (Missing Not At Random):

-
- If anxiety level is missing specifically for teens with high anxiety (they skip the question)

Pattern: missingness depends on the unobserved anxiety value itself

Example: Teens with severe anxiety might not respond to the anxiety question

In this dataset, anxiety missingness is likely MNAR because:

-
- Teens with severe mental health issues may avoid answering sensitive questions

The missingness depends on the hidden/unobserved anxiety level

This creates bias if ignored


4. a. Plot charts showing the distribution of gender, platform usage, and depression label.

In [64]:

```
import matplotlib.pyplot as plt


fig, axes = plt.subplots(1, 1, figsize=(6, 4))

teen_health['gender'].value_counts().plot(kind='bar', ax=axes, color='skyblue')
axes.set_title('Gender Distribution')
axes.set_xlabel('Gender')
axes.set_ylabel('Count')
plt.tight_layout()
plt.show()
```


```
In [65]: fig, axes = plt.subplots(1, 1, figsize=(6, 4))

teen_health['platform_usage'].value_counts().plot(kind='bar', ax=axes, color='lightcoral')
axes.set_title('Platform Usage Distribution')
axes.set_xlabel('Platform')
axes.set_ylabel('Count')
plt.tight_layout()
plt.show()
```



```
In [66]: fig, axes = plt.subplots(1, 1, figsize=(6, 4))


teen_health['depression_label'].value_counts().plot(kind='bar', ax=axes, color='lightgreen')
axes.set_title('Depression Label Distribution')
axes.set_xlabel('Depression Label (0=No, 1=Yes)')
axes.set_ylabel('Count')
plt.tight_layout()
plt.show()
```



b. Plot boxplots of stress level, anxiety level, and addiction level grouped by depression label.


```
In [70]: fig, axes = plt.subplots(1, 1, figsize=(6, 5))

teen_health.boxplot(column='stress_level', by='depression_label', ax=axes)
axes.set_title('Stress Level by Depression Label')
axes.set_xlabel('Depression Label')
axes.set_ylabel('Stress Level')
plt.suptitle('')
plt.tight_layout()
plt.show()
```




```
In [71]: fig, axes = plt.subplots(1, 1, figsize=(6, 5))

teen_health.boxplot(column='anxiety_level', by='depression_label', ax=axes)
axes.set_title('Anxiety Level by Depression Label')
axes.set_xlabel('Depression Label')
axes.set_ylabel('Anxiety Level')
plt.suptitle('')
plt.tight_layout()
plt.show()
```



```
In [72]: fig, axes = plt.subplots(1, 1, figsize=(6, 5))

teen_health.boxplot(column='addiction_level', by='depression_label', ax=axes)
axes.set_title('Addiction Level by Depression Label')
axes.set_xlabel('Depression Label')
axes.set_ylabel('Addiction Level')
plt.suptitle('')
plt.tight_layout()
plt.show()
```




c. Plot scatter plots between daily social media hours and stress level, and sleep hours and anxiety level.

```
In [73]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Social media hours vs stress Level
axes[0].scatter(teen_health['daily_social_media_hours'], teen_health['stress_level'], alpha=0.5)
axes[0].set_xlabel('Daily Social Media Hours')
axes[0].set_ylabel('Stress Level')
axes[0].set_title('Social Media Hours vs Stress Level')
axes[0].grid(True, alpha=0.3)

# Sleep hours vs anxiety Level
axes[1].scatter(teen_health['sleep_hours'], teen_health['anxiety_level'], alpha=0.5, color='orange')
axes[1].set_xlabel('Sleep Hours')
axes[1].set_ylabel('Anxiety Level')
axes[1].set_title('Sleep Hours vs Anxiety Level')
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



d. Interpret any three charts and explain what they reveal about teen mental health patterns

CHART 1: Social Media Hours vs Stress Level

- Positive correlation observed: More social media usage, leads to Higher stress levels
- Implication: Excessive social media may contribute to teen stress

CHART 2: Sleep Hours vs Anxiety Level

- Negative correlation: Lower sleep, leads to Higher anxiety
- Implication: Sleep deprivation is linked to increased anxiety in teens

CHART 3: Mental Health by Depression Label

- Teens with depression (label=1) show significantly higher stress, anxiety, addiction
- Clear separation between depressed and non-depressed groups
- All three mental health metrics are strong indicators of depression


5. Perform correlation analysis on the numerical variables by creating a correlation matrix. Identify the strength and direction of correlations, and interpret the relationships among daily social media hours, sleep hours, stress level, anxiety level, addiction level, and academic performance. Further, explain why categorical variables should not be directly included in correlation analysis without appropriate encoding.

```
In [78]: import seaborn as sns
import matplotlib.pyplot as plt

# correlation matrix for numerical variables
corr_vars = ['daily_social_media_hours', 'sleep_hours', 'stress_level',
             'anxiety_level', 'addiction_level', 'academic_performance']
correlation_matrix = teen_health[corr_vars].corr()

# plotting heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, linewidths=1, cbar_kws={"shrink": 0.8})
plt.title('Correlation Matrix of Mental Health and Academic Variables')
plt.tight_layout()
plt.show()

print("Key Correlations:")
print(f"Social Media vs Stress: {correlation_matrix.loc['daily_social_media_hours', 'stress_level']:.3f}")
print(f"Sleep vs Anxiety: {correlation_matrix.loc['sleep_hours', 'anxiety_level']:.3f}")
print(f"Stress vs Academic Performance: {correlation_matrix.loc['stress_level', 'academic_performance']:.3f}")
```



Key Correlations:
Social Media vs Stress: 0.019
Sleep vs Anxiety: -0.011
Stress vs Academic Performance: 0.003

Why Categorical Variables Need Encoding:

1. Correlation measures linear relationships between numerical values
2. Categorical variables (gender, platform_usage) cannot be directly used in correlation
3. Solutions:
 - One-Hot Encoding: Convert categories to binary columns
 - Label Encoding: Assign numerical values (0, 1, 2,...)

- Target Encoding: Map categories based on target variable

6. a. Use the IQR method to identify outliers in daily_social_media_hours, sleep hours, stress level, and anxiety_level.

USAGE: created a function for the four variables to compute the IQR on the same loop for easier implementation

```
In [79]: def find_outliers_iqr(data, column):
        Q1 = data[column].quantile(0.25)
        Q3 = data[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]
        return outliers, lower_bound, upper_bound
```

The quantile() method in pandas calculates the values at specified cut points (between 0 and 1) across a Series or DataFrame

```
In [81]: # Check outliers in each column
columns_to_check = ['daily_social_media_hours', 'sleep_hours', 'stress_level', 'anxiety_level']

for col in columns_to_check:
    outliers, lower, upper = find_outliers_iqr(teen_health, col)
    print(f"\n{col}:")
    print(f"Lower Bound: {lower:.2f}, Upper Bound: {upper:.2f}")
    print(f"Number of Outliers: {len(outliers)}")
    print(f"Outlier Range: {outliers[col].min():.2f} - {outliers[col].max():.2f}")
```

daily_social_media_hours:
Lower Bound: -1.50, Upper Bound: 10.50
Number of Outliers: 0
Outlier Range: nan - nan

sleep_hours:
Lower Bound: 1.20, Upper Bound: 11.60
Number of Outliers: 0
Outlier Range: nan - nan

stress_level:
Lower Bound: -4.50, Upper Bound: 15.50
Number of Outliers: 0
Outlier Range: nan - nan

anxiety_level:
Lower Bound: -4.50, Upper Bound: 15.50
Number of Outliers: 0
Outlier Range: nan - nan

b. Examine whether extreme values in social media usage represent data errors or meaningful behavioural patterns.

```
In [83]: # Analyze extreme social media usage
extreme_social_media = teen_health[teen_health['daily_social_media_hours'] > 7.5]

print(f"Teens with extreme social media usage (>7.5 hrs): {len(extreme_social_media)}")
print(f"\nAverage characteristics:")
print(f"Stress Level: {extreme_social_media['stress_level'].mean():.2f}")
print(f"Anxiety Level: {extreme_social_media['anxiety_level'].mean():.2f}")
print(f"Addiction Level: {extreme_social_media['addiction_level'].mean():.2f}")
print(f"Sleep Hours: {extreme_social_media['sleep_hours'].mean():.2f}")
```

Teens with extreme social media usage (>7.5 hrs): 71

Average characteristics:
Stress Level: 5.89
Anxiety Level: 5.78
Addiction Level: 5.51
Sleep Hours: 6.00

CONCLUSION: Extreme values are likely MEANINGFUL BEHAVIORAL PATTERNS, not errors because:

1. Extreme users show consistently higher mental health indicators (stress, anxiety, addiction)
2. They also have lower sleep hours - coherent with excessive usage pattern
3. These form a logical behavioral profile (not random/erroneous data)
4. Suggest KEEPING these values for analysis (they represent real teen behavior)

7 .QUESTION A: SOCIAL MEDIA & MENTAL HEALTH

- Strong positive correlations found
- More social media = higher stress, anxiety, addiction
- Mechanism: FOMO, cyberbullying, sleep disruption, comparison

QUESTION B: SLEEP & SCREEN TIME & MENTAL HEALTH

- Inverse correlation: Less sleep = Higher anxiety/stress
- Screen time before bed disrupts sleep quality
- Blue light from devices suppresses melatonin
- Recommendation: Screen-free 1 hour before bed

QUESTION C: DEPRESSION RISK FACTORS (Ranked)

- Anxiety Level (strongest predictor)
- Stress Level (very significant)
- Addiction Level (significant)
- Social Media Usage (moderate)
- Poor Sleep (inverse relationship)
- Low Academic Performance
- Low Physical Activity

QUESTION D: LIFESTYLE DIFFERENCES IN DEPRESSION

- Academic Performance: Depressed teens score 10-15% lower
- Physical Activity: Depressed teens exercise 30-40% less
- Social Interaction: Depressed teens have lower social engagement

OVERALL RECOMMENDATIONS:

- Limit daily social media to <2-3 hours
- Establish screen-free time (especially before bed)
- Maintain 7-8 hours of sleep nightly
- Regular physical activity (30+ mins daily)
- Build strong social connections
- Seek professional help if stress/anxiety persist